

A Machine Learning Pipeline for Geographic Source Attribution of Shiga Toxin–Producing *Escherichia coli* Using Random Forest Classification

Guillermo Comesaña Cimadevila^{1*} and Lauren Cowley¹

¹Department of Life Sciences, University of Bath, Claverton Down, Bath BA2 7AY, UK

*Corresponding author: gcc46@bath.ac.uk

1 Introduction

Advances in pathogen genomics have transformed public health microbiology, enabling high-resolution identification, surveillance, and source attribution of infectious diseases [1,2]. Central to this transformation is whole-genome sequencing (WGS), which offers a comprehensive view of microbial diversity, evolutionary dynamics, and transmission pathways at an unprecedented scale [3,4]. However, the exponential growth in sequencing throughput has exposed the limitations of traditional bioinformatics pipelines, particularly for predictive tasks that rely on complex, polygenic genomic signals—such as antimicrobial resistance (AMR) prediction and geographic source attribution [4,5].

Machine learning (ML) has emerged as a scalable approach to these challenges, offering the ability to model high-dimensional, non-linear relationships inherent in genomic data [1,2]. In AMR prediction, ML algorithms trained on k-mer frequencies, single nucleotide polymorphisms (SNPs), and gene presence-absence matrices have achieved high predictive performance, with accuracies exceeding 95% for pathogens such as *Escherichia coli* and *Mycobacterium tuberculosis* [6,7]. These models underpin diagnostic platforms such as Mykrobe, ResFinder, and DeepARG, enabling rapid resistance profiling and guiding targeted antimicrobial interventions [8,9,10].

By contrast, the application of ML to geographic source attribution remains comparatively underexplored, despite its critical role in outbreak investigations, targeted surveillance, and population-specific public health responses [3,11]. Recent studies employing multivariate and hierarchical classification frameworks have reported promising results, with macro-F1 scores up to 0.95 at the continental level and 0.66 at the national level [3,11]. Nevertheless, these models often exhibit limited generalisability across species with diverse population structures and reduced reproducibility on independent datasets. Additional challenges—including high computational demands, inconsistent feature representations, and limited interpretability—continue to hinder their integration into real-time public health workflows [4].

A key distinction between AMR prediction and geographic attribution underlies many of these challenges. AMR is often mediated by discrete, well-characterised genetic elements such as resistance genes (e.g., *bla*-CTX-M, *mecA*) and SNPs (e.g., mutations in *gyrA* conferring fluoroquinolone resistance), which are often associated with strong phenotypic effects [12]. Geographic signals, by contrast, are weak, polygenic, and diffusely distributed across the genome. They arise from complex

processes such as recombination, horizontal gene transfer, population stratification, and regional selective pressures, and are further confounded by uneven sampling densities and incomplete geographic coverage [4]. Consequently, geographic attribution requires models capable of detecting subtle, genome-wide patterns that are not captured by canonical genetic markers.

Phylogenetic reconstruction remains the main method for inferring geographic origin, offering robust evolutionary inference based on shared ancestry [3,11]. However, these approaches are computationally intensive, require specialised expertise, and scale poorly with modern genomic surveillance data [13]. ML-based alternatives offer a complementary strategy, enabling the direct inference of geographic origin from raw genomic data without the need for sequence alignment or phylogenetic reconstruction. Such approaches promise substantial gains in speed, scalability, and automation, making them well-suited to public health settings [1,2].

Shiga toxin-producing *Escherichia coli* (STEC) is a globally distributed zoonotic pathogen and a priority target for better attribution methods. STEC causes a spectrum of clinical outcomes, from mild diarrhoea to haemorrhagic colitis and haemolytic uraemic syndrome, and remains a leading cause of foodborne illness and hospitalisation [14]. While its virulence factors and outbreak dynamics are well characterised, geographic attribution of STEC isolates remains hindered by fragmented datasets, inconsistent analytical frameworks, and the absence of reproducible, scalable pipelines. In 2021, the UK reported the highest STEC notification rate in Europe [15], underscoring the urgency of improving attribution accuracy to support national surveillance and guide public health interventions.

In this study, we benchmark a supervised ML approach for the geographic attribution of STEC isolates using k-mer-derived genomic features. We train and evaluate Random Forest classifiers for country- and region-level prediction, focusing on robustness to population structure, reproducibility across held-out test data, and feature-level interpretability. Our results highlight both the promise and current limitations of k-mer-based ML in this context, and provide a reproducible pipeline for further methodological development in genomic epidemiology.

2 Methods

2.1 Data Acquisition and Metadata Harmonisation

A schematic overview of the pipeline is presented in Figure 1. WGS data for STEC isolates were retrieved from the National Center for Biotechnology Information (NCBI) Sequence Read Archive (SRA; <https://www.ncbi.nlm.nih.gov/sra>). The training set consisted of 2,459 isolates collected between 2014 and 2018, while an independent test set comprising 415 isolates from 2019 was held out for final evaluation. Associated metadata were obtained in CSV format and included BioSample accession numbers, country and region of origin, Shiga toxin (Stx) subtype, phage type (PT), and a precomputed k-mer count matrix for each isolate.

Metadata harmonisation involved structured cleaning. All text fields were normalised by converting to lowercase, trimming whitespace, and removing punctuation. Country names were standardised using a controlled vocabulary—for example, “U.S.A.”, “USA”, and “United States of America” were unified under “United States”—and constituent nations of the UK (e.g., England, Wales) were grouped under “UK”. Duplicate BioSample accession numbers were resolved by retaining the first valid occurrence, and samples missing or ambiguously annotated for Stx or PT (i.e. “unknown”, “NaN”) were excluded. Labels appearing in fewer than ten training samples were grouped under “Other Country” and “Other Region” to reduce class fragmentation, consistent with standard classification practices [16]. To preserve label alignment, any test sample containing a country or region label absent in the training data was excluded. Following harmonisation and alignment with k-mer data, 2,335 isolates were retained in the training set, and 398 isolates in the test set, suitable for both country- and region-level classification.

2.2 K-mer Feature Matrix Construction and Normalisation

Precomputed k-mer count matrices were obtained from NCBI, generated using the iMOKA platform [17], which employs the KMC3 algorithm [18] to extract canonical 31-base k-mers from raw sequencing reads. Each isolate was represented as a sparse vector of k-mer counts. Matrices were transposed to match samples (rows) and features (columns). To reduce sparsity, k-mers present in fewer than 5% of training samples were removed, a standard threshold for eliminating noise

due to sequencing artefacts and low-frequency mutations in k-mer-based classification [19]. This reduced the dimensionality from 49,999 to 30,070 features. Raw counts were converted to relative abundances by dividing by the total k-mer count per sample and multiplying by 100. The matrix was converted to 64-bit and z-score standardised.

2.3 Feature Selection

Feature selection was conducted in three sequential steps, using only the training data to prevent data leakage. First, features with near-zero variance ($\sigma^2 < 1 \times 10^{-4}$) were removed to discard uninformative predictors. This threshold eliminates k-mers with negligible variation across samples, which contribute little to classification and may introduce noise [20]. Second, pairwise Pearson correlation coefficients were computed among the remaining features. For any pair with $|r| < 0.95$, one feature was deterministically removed to mitigate multicollinearity, which is common among overlapping k-mers and can distort feature importance estimates [21].

Third, a Random Forest classifier was trained on the unbalanced training set to rank features by mean decrease in Gini impurity. This step was intentionally performed prior to class balancing, ensuring that feature selection captured true distributional patterns, avoiding bias from synthetic re-sampling [22]. The Random Forest used 500 estimators, a maximum tree depth of 20, and class-weighted subsampling (`class_weight='balanced_subsample'`), consistent with prior k-mer-based approaches [23,24]. Features were ranked in descending order of importance, and the top 2,000 were retained. This cap was selected to balance predictive power with computational efficiency, as performance gains plateau beyond this range [25]. The final feature sets contained 2,000 features for the region task and 1,857 for the country task, based on their respective rank distributions.

2.4 Class Balancing

The training data exhibited extreme class imbalance, with the “UK” class comprising 77.4% of samples at the country level and 77.3% at the region level. To address this, three resampling strategies were evaluated: random oversampling, random undersampling, and Synthetic Minority Oversampling Technique (SMOTE) [26]. Balance was assessed using three standard metrics: Gini

index (G), normalised Shannon entropy (H_{norm}), and the imbalance ratio (IR), defined respectively as:

$$G = 1 - \sum_{i=1}^C p_i^2 \quad (1)$$

$$H_{\text{norm}} = \frac{-1}{\log_2 C} \sum_{i=1}^C p_i \log_2 p_i \quad (2)$$

$$IR = \frac{\max(p_i)}{\min(p_i)} \quad (3)$$

where C is the number of classes and p_i is the relative frequency of class i . Prior to resampling, $G = 0.3971$, $H_{\text{norm}} = 0.3002$, and $IR = 1,807$ for country, and $G = 0.3913$, $H_{\text{norm}} = 0.3857$, and $IR = 604$ for region. Each method succeeded in achieving perfect balance ($G = 0.9853$, $H_{\text{norm}} = 1.0$, and $IR = 1.0$), but SMOTE consistently yielded higher macro F1 scores across five-fold stratified cross-validation. SMOTE was therefore selected, and all minority classes were oversampled to match the majority class ("UK"), resulting in 1,807 samples per class.

2.5 Model Training and Evaluation

To optimise the training-validation split, three stratified ratios commonly used in pathogen geographic source attribution (70:30, 75:25, and 80:20) were evaluated on the SMOTE-balanced training set. The 80:20 split consistently produced the highest macro F1 score and was therefore selected. Final Random Forest classifiers were trained with `class_weight='balanced'` to address any residual imbalance.

Hyperparameter tuning was performed using GridSearchCV over 120 configurations, evaluating five-fold stratified cross-validation and optimising for macro F1. Parameters included: number of estimators (1–200), tree depth (1–40), number of features per split (sqrt, log2, 0.1, 0.25, 0.5, 0.75, 1.0), minimum samples to split (2, 5, 10), and minimum samples per leaf (1, 2, 5). The optimal configuration was: 200 estimators, depth 20, sqrt for feature selection, 2 samples to split, and 1

sample per leaf. This model was retrained on the full training set and evaluated on both internal validation and the 2019 test set.

Performance was measured using per-class and aggregate precision, recall, macro- and weighted F1 scores, and accuracy. Confusion matrices were generated to visualise misclassification. Feature importance was computed via mean decrease in Gini impurity, and the top ten features were extracted for biological interpretation.

2.6 K-mer Annotation

The top 10 features (k-mers) for each task were annotated using BLASTn against the NCBI nucleotide database (<https://www.ncbi.nlm.nih.gov/nucleotide/>). Queries were constructed in FASTA format and run using BLAST+ v2.13.0 [27] with remote querying enabled (-remote), tabular output format (-outfmt 6), and one target alignment per query (-max_target_seqs 1). Default scoring parameters were used, including a word size of 11, match and mismatch scores of +2 and 3 respectively, and gap penalties of 5 for opening and 2 for extension. A reproducible shell script (scripts/03_blast_top10.sh) integrated into the pipeline handled all BLAST calls. Returned alignments were analysed to assess whether the top k-mers matched known functional elements such as stx genes, O-antigen loci, or phage insertion sites. These mappings validated classifier output and clarified geographic variation in pathogenicity.

2.7 Computational Reproducibility and Code Availability

All analyses were conducted in a Conda-managed Python 3.12 environment. Key libraries included scikit-learn (1.4.1), numpy (1.26.4), pandas (2.2.2), seaborn (0.13.2), imbalanced-learn (0.12.1), and joblib (1.4.0). A global random seed (42) was applied to all stochastic processes. The pipeline was orchestrated with Nextflow (v22.10.1) and containerised using Docker for full reproducibility. All experiments were executed on an AWS r5n.8xlarge instance (32 vCPUs, 256 GiB RAM) running Ubuntu 24.04 in the eu-west-2 region. Code, data, models, and documentation are publicly available at: <https://github.com/guillermocomesanacimadevila/STEC-Classififer>.




Figure 1: Schematic Overview of the STEC Random Forest Pipeline

3 Results




Figure 2: **Class distributions across training and test sets for region and country classification tasks.** Panels show the number of samples per class at two geographic levels: Region (A) and Country (B). Each bar represents the number of samples assigned to a given class across three datasets: the training set before SMOTE ("Pre-SMOTE"), the training set after SMOTE ("Post-SMOTE"), and the 2019 test set ("Test Set").

Table 1: Class distribution and diversity metrics for the 2019 test set.

Task	# Classes	Total Samples	Avg./Class	Gini Index	Shannon Entropy	Imbalance Ratio
Region	11	398	36.2	0.56	0.54	254.0
Country	15	398	26.5	0.57	0.51	254.0

Note: Summary statistics for region- and country-level classification tasks in the 2019 independent test set, including class count, average samples per class, and three diversity metrics.

The dataset comprised 2,733 STEC isolates, with 2,335 collected between 2014 and 2018 forming the training set, and 398 from 2019 reserved as an independent test set. Isolates were annotated at both country and regional levels, spanning 15 countries and 11 regions. In the 2019 test set, 77.4% of samples originated from the UK, resulting in substantial class imbalance. For region-level classification, this corresponded to $\mathbf{G} = 0.56$, $H_{\text{norm}} = 0.54$, and $\text{IR} = 254.0$. The country-level task showed comparable skew, with $\mathbf{G} = 0.57$, $H_{\text{norm}} = 0.51$, and $\text{IR} = 254.0$. The original training

data exhibited a similar imbalance, with over 77% of isolates again assigned to the UK. To correct for this, the training set was balanced using SMOTE, increasing minority class sizes to match the UK’s sample count. This produced uniform class distributions ($\mathbf{G} = 0.9853$, $H_{\text{norm}} = 1.0$, $\text{IR} = 1.0$) across both classification tasks. Class distributions before and after balancing are shown in Figure 2. The test set remained unchanged to preserve a realistic distribution for final evaluation. Random Forest classifiers trained on the balanced data were evaluated on the 2019 test set. For region-level prediction, the model achieved an accuracy of 0.827, with macro F1 = 0.572, weighted F1 = 0.855, precision = 0.514, and recall = 0.740. For country-level classification, accuracy reached 0.701, with macro F1 = 0.419, weighted F1 = 0.751, precision = 0.365, and recall = 0.638 (see Table 2).

Table 2: Performance metrics for region and country classification on the 2019 test set.

Task	Accuracy	Macro F1	Weighted F1	Precision / Recall
Region Classification	0.827	0.572	0.855	0.514 / 0.740
Country Classification	0.701	0.419	0.751	0.365 / 0.638

Note: Accuracy, F1 scores, and precision/recall values are reported for each task. Macro scores treat all classes equally, while weighted scores reflect class support.




Figure 3: **Confusion matrices for region and country classification on the 2019 test set.** (A) Confusion matrix for region classification; (B) Confusion matrix for country classification. Rows correspond to the true class labels, and columns to predicted labels. Cell values represent the number of samples assigned to each predicted class. Color intensity reflects prediction count. Only test set data were used for evaluation.

Performance differed substantially across classes. In the region task, F1 scores were highest for the UK (0.900), the Middle East (0.894), and Southern Europe (0.844). Moderate scores were observed for Central Europe, North Africa, and Northern Europe. Sub-Saharan Africa returned an F1 of 0.000, and Other Region scored 0.118. The confusion matrix in Figure 3A shows strong separation for dominant classes and increased misclassification among those with low representation. At the country level, the best performance was seen for Portugal (F1 = 0.889), Turkey (0.873), and Egypt (0.727), followed by Cyprus, Spain, and the UK. Countries including India, Ireland, and Mexico received F1 scores of 0.000, reflecting total misclassification. The confusion matrix in Figure 3B highlights broader dispersion in predictions for this task, particularly among minority classes.




Figure 4: Top 10 k-mers for region and country classification ranked by Gini importance. (A) Highest-ranked k-mers in the region-level Random Forest model. (B) Top features for country-level classification. Each bar represents a k-mer feature, with Gini importance reflecting its relative contribution to the model. Labels are truncated for clarity where necessary.

To identify the most informative genomic features, Gini importance scores were computed across the final feature sets of 2,000 k-mers for region classification and 1,857 for country. The region model's top ten features ranged from 0.0241 to 0.0083 in importance, while those for the country model ranged from 0.0191 to 0.0129, as shown in Figure 4. BLASTn analysis of these top-ranked k-mers revealed consistent alignment to *E. coli* O157:H7 genomes, frequently mapping to phage-associated or accessory regions. None matched to known virulence genes, and several k-mers of 11 nucleotides or fewer failed to produce high-confidence alignments. Representative matches are reported in Tables 3 and 4.

Table 3: Representative BLAST hits for the top-ranked k-mers in region classification.

Rank	K-mer (truncated)	Importance	Matched Gene	Species	E-value	% Identity
1	TTTTTTTGCT...TTTTT	0.0241	—	<i>E. coli O157:H7</i>	7e-43	100.0%
2	TTTTTGTGGC...AGGGG	0.0235	—	<i>E. coli O157:H7</i>	6e-11	100.0%
3	TTTTTTAACG...AGGAA	0.0151	—	<i>E. coli O157:H7</i>	7e-43	100.0%
4	TTTTTGAGTT...TTGTG	0.0142	—	<i>E. coli O157:H7</i>	7e-43	100.0%
5	TTTTTGAATC...GTTTCG	0.0141	—	<i>E. coli strain 164</i>	7e-43	100.0%
6	TTTTTGCAACAGC	0.0105	—	—	—	—
7	TTTTTTCCGC...TTTGT	0.0104	—	<i>E. coli strain 155</i>	7e-43	100.0%
8	TTTTTGCTATCC	0.0094	—	—	—	—
9	TTTTTTGACGGG	0.0087	—	—	—	—
10	TTTTTGATCGCTGA	0.0083	—	—	—	—

Note: K-mer sequences are truncated for readability. Full-length versions are in Supplementary Table 9 (S5A). Columns report Gini importance, matched gene (if any), species, E-value, and percent identity. Dashes (—) indicate no confident match.

Table 4: Representative BLAST hits for the top-ranked k-mers in country classification.

Rank	K-mer (truncated)	Importance	Matched Gene	Species	E-value	% Identity
1	TTTTTTCTAT...GGTAC	0.0191	—	<i>E. coli O157:H7</i>	2e-44	100.0%
2	TTTTTTGTGG...ATAAG	0.0179	—	<i>E. coli O157:H7 strain Z1815</i>	1e-08	100.0%
3	TTTTTGTTTCG...AGGTG	0.0173	—	<i>E. coli O157:H7 FLT_1998F</i>	1e-46	100.0%
4	TTTTTGCCGC...GCCAG	0.0153	—	<i>E. coli O157:H7 strain Z1815</i>	1e-46	100.0%
5	TTTTTGCTATG	0.0142	—	—	—	—
6	TTTTTGTTAA...TATCA	0.0140	—	<i>E. coli O157:H7 FLT_2000A</i>	1e-46	100.0%
7	TTTTTTGTCC...CATTT	0.0135	—	<i>E. coli O157:H7 strain Z1815</i>	1e-46	100.0%
8	TTTTTGTTTTTTCT	0.0133	—	—	—	—
9	TTTTTTCTCGCTC	0.0130	—	—	—	—
10	TTTTTTCCGCTGG	0.0129	—	—	—	—

Note: Truncated k-mers shown. Full sequences in Supplementary Table 10 (S5B). Columns as in Table 3.

On the SMOTE-balanced training set, both classifiers achieved near-perfect performance, with accuracy, macro F1, and weighted F1 all exceeding 0.99. Confusion matrices for the training data (Supplementary Figure 7) showed clean diagonal alignment and minimal error. Complete per-class precision, recall, and F1 scores are reported in Supplementary Tables 5 through 8. Hyperparameter tuning curves and selected configurations for both models are presented in Supplementary Figures 5 and 6.

4 Discussion

4.1 Model Performance and Biological Interpretation

This study demonstrates the potential and limitations of ML in attributing the geographic origin of STEC using k-mer-based genomic features. At the regional level, the Random Forest classifier achieved moderate predictive performance (macro F1 = 0.572), suggesting that diffuse phylogeographic signals are detectable within the accessory genome. However, performance deteriorated significantly at the country level (macro F1 = 0.419), highlighting the challenge of resolving fine-scale population structure using current feature representations and classification strategies.

The most significant constraint was severe class imbalance. Over 77% of the dataset originated from the UK, while most other countries were underrepresented. Although SMOTE was applied to generate a balanced training set, its underlying assumptions—local linearity and continuity in Euclidean space—are incompatible with the evolutionary dynamics of microbial genomes [28]. Processes such as homologous recombination, horizontal gene transfer, and selective sweeps introduce non-linear variation that cannot be accurately captured through synthetic interpolation [29].

Additionally, while the implementation of an 'Other' category improved class balance metrics and reduced label fragmentation during training, it introduces a critical trade-off: classes merged into 'Other' cannot be individually resolved. As a result, any future samples belonging to those original minority classes will be systematically misclassified, undermining the model's utility for fine-grained geographic attribution and limiting generalisability to underrepresented or novel sources. Taken together, these balancing strategies—synthetic oversampling and label collapsing—reduced biological realism and granularity in the training data. Consequently, SMOTE-generated samples replicated existing patterns without introducing evolutionarily meaningful variation. As a result, the model learned artificial correlations, overfit to synthetic data, and failed to generalise to real genomic sequences, indicating that such models are unlikely to transfer effectively to independent datasets with different population structures or sampling distributions [30].

This overfitting was evident in the disparity between training and test performance. While macro F1 scores exceeded 0.99 on the balanced training data, predictive accuracy declined substantially on the 2019 test set. The model learned artefactual patterns embedded in synthetic samples rather than

robust genomic features reflective of geographic origin [31]. Confusion matrices revealed a strong predictive bias toward the dominant UK class, with minority-class isolates frequently misclassified, indicating that the classifier captured general population structure but failed to isolate class-specific genomic signals.

Feature importance analysis further underscored these limitations. The top-ranked k-mers, identified by Gini importance, predominantly mapped to phage-associated regions, intergenic loci, and hypothetical proteins—components of the accessory genome known for their rapid turnover and limited phylogenetic stability [32]. Notably, none of the most informative k-mers aligned to conserved core genes or established virulence determinants such as *stx*, nor did they produce high-confidence BLAST alignments in many cases. This suggests that the model relied on unannotated or artefactual genomic variation, limiting both biological interpretability and practical applicability in public health contexts [1].

4.2 Contextualisation within the Literature

The observed generalisation gap aligns with findings from previous microbial attribution studies. Pascoe et al [33] reported that a Random Forest model trained on over 25,000 *Campylobacter jejuni* genomes achieved 76% accuracy in internal validation but only 49% on external datasets, illustrating a common challenge in modelling population-structured pathogens [33]. In contrast, Buultjens et al. [11] achieved over 90% attribution accuracy for *Legionella pneumophila* by leveraging isolates from controlled environmental sources with structured metadata, allowing the model to learn well-defined, source-specific genomic patterns [11]. The absence of such contextual metadata in this study—such as host species, environmental exposure, or isolation source—precluded the model from anchoring genomic variation to ecologically or epidemiologically meaningful axes [34].

A key methodological difference between this study and higher-performing frameworks lies in the classification architecture. Bayliss et al. [3] employed a hierarchical ML strategy for *Salmonella enterica* attribution, which decomposed prediction into nested geographic levels. Their model achieved macro F1 scores of 0.954 at the continental level, 0.718 at the regional level, and 0.661 at the country level, reflecting a structure-aligned task formulation [3]. In contrast, the flat classification strategy adopted here required simultaneous resolution across all country-level classes, many

of which exhibited high genomic similarity and low representation. This led to greater inter-class confusion and reduced discriminatory power, particularly in regions with overlapping population structures [35].

4.3 Methodological Reflections and Future Directions

Random Forests were selected for their interpretability and computational efficiency, providing a useful baseline for genomic classification. However, Random Forests are limited in their ability to capture complex interactions between features and to optimise decision boundaries globally across the feature space [36]. Gradient-boosted frameworks such as XGBoost and LightGBM may offer enhanced performance, particularly under imbalanced conditions, by focusing on hard-to-classify instances through iterative reweighting [37]. Similarly, kernel-based methods—such as support vector machines with locality-preserving dimensionality reduction—may improve classification in high-dimensional, sparse feature spaces typical of k-mer matrices [38].

The reliance on Gini importance for feature selection introduced additional bias. This metric disproportionately favours frequent features and fails to capture rare, lineage-specific variants that may be highly informative in microbial source attribution [39,40]. Model-agnostic tools such as SHAP (Shapley Additive Explanations) could improve interpretability by attributing prediction contributions at the per-sample level, enabling identification of discriminative features even in genetically heterogeneous populations [41].

Although deep learning approaches were not implemented in this pipeline, they represent a promising direction for future work. Convolutional neural networks (CNNs) and transformer-based models can learn sequence-level features directly from raw genomic data, potentially capturing both local motifs and broader structural signals [1]. However, microbial datasets are often sparse and weakly labelled, posing challenges for supervised training [41]. Emerging techniques such as transfer learning and self-supervised pretraining may mitigate these limitations and enable deep models to learn biologically meaningful representations even in low-resource settings [42]. Ultimately, the integration of structured metadata—including host, geography, time, and environment—will be essential to disambiguate weak genomic signals and ground them in biologically interpretable axes.

4.4 Conclusions

In conclusion, this study demonstrates that k-mer-based ML models can identify broad phylogeographic signals in STEC genomes but lack sufficient resolution to generalise at finer scales. Class imbalance, biologically implausible oversampling, and a flat classification strategy collectively undermined model performance and interpretability. The reliance on unannotated accessory elements as primary features further limited the biological relevance of the results. Advancing microbial source attribution will require methods that align with microbial population structure, models capable of handling high-dimensional sparsity and class imbalance, and interpretability frameworks that identify meaningful genomic features. Critically, the availability and incorporation of structured metadata will be pivotal in transforming statistical predictions to contribute meaningfully to public health surveillance [24,43].

References

1. Asnicar F, Thomas AM, Passerini A, Waldron L, Segata N. Machine learning for microbiologists. *Nature Reviews Microbiology* [Internet]. 2023 Nov 15 [cited 2023 Nov 23];1–15. Available from: https://www.nature.com/articles/s41579-023-00984-1?fbclid=IwAR22MDMA_9Fk5TCBG9oPgZ9QyJpfstWb7DgvU75BIxtr6p5JiD944sIXWgE
2. Gao Y, Liu M. Application of machine learning based genome sequence analysis in pathogen identification. *Frontiers in Microbiology*. 2024 Oct 2;15.
3. Bayliss SC, Locke RK, Jenkins C, Marie Anne Chattaway, Dallman TJ, Cowley LA. Rapid geographical source attribution of *Salmonella enterica* serovar Enteritidis genomes using hierarchical machine learning. *eLife*. 2023 Apr 12;12.
4. Georgios Papoutsoglou, Tarazona S, Lopes MB, Klammersteiner T, Ibrahimi E, Eckenberger J, et al. Machine learning approaches in microbiome research: challenges and best practices. *Frontiers in Microbiology* [Internet]. 2023 Sep 22 [cited 2024 Jan 26];14. Available from: <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC10556866/>
5. Angers-Loustau A, Petrillo M, Bengtsson-Palme J, Berendonk T, Blais B, Chan KG, et al. The challenges of designing a benchmark strategy for bioinformatics pipelines in the identification of antimicrobial resistance determinants using next generation sequencing technologies. *F1000Research*. 2018 Dec 7;7:459.
6. Ali T, Ahmed S, Aslam M. Artificial Intelligence for Antimicrobial Resistance Prediction: Challenges and Opportunities towards Practical Implementation. *Antibiotics*. 2023 Mar 6;12(3):523.
7. Abhadionmhen AO, Asogwa CN, Ezema ME, Nzeh RC, Ezeora NJ, Abhadiomhen SE, et al. Machine Learning Approaches for Microorganism Identification, Virulence Assessment, and Antimicrobial Susceptibility Evaluation Using DNA Sequencing Methods: A Systematic Review. *Molecular biotechnology* [Internet]. 2024 Sep; Available from: <https://pubmed.ncbi.nlm.nih.gov/39520638/>

-
8. Florensa AF, Kaas RS, Clausen PTLC, Aytan-Aktug D, Aarestrup FM. ResFinder – an open online resource for identification of antimicrobial resistance genes in next-generation sequencing data and prediction of phenotypes from genotypes. *Microbial Genomics* [Internet]. 2022 Jan 24;8(1). Available from: <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC8914360/pdf/mgen-8-0748.pdf>
 9. Hunt M, Bradley P, Lapierre SG, Heys S, Thomsit M, Hall MB, et al. Antibiotic resistance prediction for *Mycobacterium tuberculosis* from genome sequence data with Mykrobe. *Wellcome Open Research* [Internet]. 2019;4:191. Available from: <https://pubmed.ncbi.nlm.nih.gov/32055708/>
 10. Arango-Argoty G, Garner E, Pruden A, Heath LS, Vikesland P, Zhang L. DeepARG: a deep learning approach for predicting antibiotic resistance genes from metagenomic data. *Microbiome*. 2018 Feb 1;6(1).
 11. Buultjens AH, Koen Vandelannoote, Karolina Mercoulia, Ballard S, Sloggett C, Howden BP, et al. High performance *Legionella pneumophila* source attribution using genomics-based machine learning classification. *Applied and Environmental Microbiology*. 2024 Jan 30;
 12. Jian Z, Zeng L, Xu T, Sun S, Yan S, Yang L, et al. Antibiotic resistance genes in bacteria: Occurrence, spread, and control. *Journal of Basic Microbiology*. 2021 Oct 14;61(12):1049–70.
 13. Chen Z, Philippe Lemey, Yu H. Approaches and challenges to inferring the geographical source of infectious disease outbreaks using genomic data. *The Lancet Microbe* [Internet]. 2023 Nov 30;5(1):e81–92. Available from: [https://www.thelancet.com/journals/lanmic/article/PIIS2666-5247\(23\)00296-3/fulltext](https://www.thelancet.com/journals/lanmic/article/PIIS2666-5247(23)00296-3/fulltext)
 14. Greig DR, Quinn OI, Rodwell EV, Olonade I, Swift C, Douglas A, et al. Genomic analysis of an outbreak of Shiga toxin-producing *Escherichia coli* O183:H18 in the United Kingdom, 2023. *Microbial genomics*. 2024 May 21;10(5).
 15. European Centre for Disease Prevention and Control. STEC infection - Annual Epidemiological Report for 2021 [Internet]. www.ecdc.europa.eu. 2022. Available from: <https://www.ecdc.europa.eu/en/publications-data/stec-infection-annual-epidemiological-report-2021>
 16. Breiman L. Random Forests. *Machine Learning* [Internet]. 2001 Oct;45(1):5–32. Available from:

<https://link.springer.com/article/10.1023/A:1010933404324>

17. Lorenzi C, Barriere S, Villemin JP, Dejardin Bretones L, Mancheron A, Ritchie W. iMOKA: k-mer based software to analyze large collections of sequencing data. *Genome biology* [Internet]. 2020 Spring;21(1):261. Available from: <https://pubmed.ncbi.nlm.nih.gov/33050927/>
18. Kokot M, Dlugosz M, Deorowicz S. KMC 3: counting and manipulating k-mer statistics. *Bioinformatics*. 2017 May 4;33(17):2759–61.
19. Arning N, Sheppard SK, Bayliss S, Clifton DA, Wilson DJ. Machine learning to predict the source of campylobacteriosis using whole genome data. Hughes D, editor. *PLOS Genetics* [Internet]. 2021 Oct 18 [cited 2025 May 7];17(10):e1009436. Available from: <https://pmc.ncbi.nlm.nih.gov/articles/PMC8553134/>
20. Deo RC. Machine Learning in Medicine. *Circulation* [Internet]. 2015 Nov 17;132(20):1920–30. Available from: <https://pubmed.ncbi.nlm.nih.gov/26572668/>
21. Mei K, Tan M, Yang Z, Shi S. Modeling of Feature Selection Based on Random Forest Algorithm and Pearson Correlation Coefficient. *Journal of Physics: Conference Series*. 2022 Apr 1;2219(1):012046.
22. Menze BH, Kelm BM, Masuch R, Himmelreich U, Bachert P, Petrich W, et al. A comparison of random forest and its Gini importance with standard chemometric methods for the feature selection and classification of spectral data. *BMC Bioinformatics*. 2009;10(1):213.
23. Orozco-Arias S, Candamil-Cortés MS, Jaimes PA, Piña JS, Reinel Tabares-Soto, Guyot R, et al. K-mer-based machine learning method to classify LTR-retrotransposons in plant genomes. *PeerJ*. 2021 May 19;9:e11456–6.
24. Tian Q, Zhang P, Zhai Y, Wang Y, Zou Q. Application and Comparison of Machine Learning and Database-Based Methods in Taxonomic Classification of High-Throughput Sequencing Data. *Genome Biology and Evolution* [Internet]. 2024 May 1;16(5). Available from: <https://academic.oup.com/gbe/article/16/5/evae102/7674165>
25. Ramos-Pérez I, Álvar Arnaiz-González, Rodríguez JM, César García-Osorio. When is resampling

beneficial for feature selection with imbalanced wide data? 2022 Feb 1;188:116015–5.

26.Chawla NV, Bowyer KW, Hall LO, Kegelmeyer WP. SMOTE: Synthetic Minority Over-sampling Technique. *Journal of Artificial Intelligence Research* [Internet]. 2002 Jun 1;16(16):321–57. Available from: <https://www.jair.org/index.php/jair/article/view/10302>

27.NCBI. BLAST: Basic Local Alignment Search Tool [Internet]. Nih.gov. 2025. Available from: <https://blast.ncbi.nlm.nih.gov/Blast.cgi>

28.Polz MF, Alm EJ, Hanage WP. Horizontal gene transfer and the evolution of bacterial and archaeal population structure. *Trends in Genetics*. 2013 Mar;29(3):170–5.

29.Liang D, Zhang H, Yuan D, Zhang M. Multi-Head Encoding for Extreme Label Classification. *IEEE Transactions on Pattern Analysis and Machine Intelligence* [Internet]. 2024 Dec 25;47(3):2199–211. Available from: <https://arxiv.org/abs/2412.10182>

30.Liu CC, William. Machine learning reveals the dynamic importance of accessory sequences for *Salmonella* outbreak clustering. *mBio*. 2025 Jan 28;

31.James G, Witten D, Hastie T, Tibshirani R. *An Introduction to Statistical Learning* [Internet]. Springer Texts in Statistics. New York, NY: Springer US; 2021. Available from: <https://link.springer.com/book/10.1007/978-1-0716-1418-1>

32.Kim JI, Maguire F, Tsang KK, Gouliouris T, Peacock SJ, McAllister TA, et al. Machine Learning for Antimicrobial Resistance Prediction: Current Practice, Limitations, and Clinical Perspective. *Clinical Microbiology Reviews*. 2022 Sep 21;35(3).

33.Pascoe B, Fitcher G, Pensar J, Bayliss SC, Mourkas E, Calland JK, et al. Machine learning to attribute the source of *Campylobacter* infections in the United States: A retrospective analysis of national surveillance data. *Journal of Infection* [Internet]. 2024 Sep 7;89(5):106265. Available from: <https://www.sciencedirect.com/science/article/pii/S0163445324001993>

34.Sudasinghe PG, Wijesinghe CR, Weerasinghe AR. Prediction of horizontal gene transfer in *Escherichia coli* using machine learning. 2013 International Conference on Advances in ICT for Emerging Regions (ICTer) [Internet]. 2013 Dec;118–24. Available from:

<https://ieeexplore.ieee.org/document/6761165>

35.Rezende PM, Xavier JS, Ascher DB, Fernandes GR, Pires DEV. Evaluating hierarchical machine learning approaches to classify biological databases. *Briefings in Bioinformatics* [Internet]. 2022 Jun 21 [cited 2025 Apr 22];23(4). Available from: <https://doi.org/10.1093/bib/bbac216>

36.Nguyen TT, Huang JZ, Nguyen TT. Unbiased Feature Selection in Learning Random Forests for High-Dimensional Data. *The Scientific World Journal* [Internet]. 2015 Mar 24 [cited 2023 Mar 10];2015:e471371. Available from: <https://www.hindawi.com/journals/tswj/2015/471371/>

37.Imani M, Beikmohammadi A, Arabnia HR. Comprehensive Analysis of Random Forest and XGBoost Performance with SMOTE, ADASYN, and GNUS Under Varying Imbalance Levels. *Technologies* [Internet]. 2025 Feb 20;13(3):88. Available from: <https://www.mdpi.com/2227-7080/13/3/88>

38.Du KL, Jiang B, Lu J, Hua J, Swamy MNS. Exploring Kernel Machines and Support Vector Machines: Principles, Techniques, and Future Directions. *Mathematics* [Internet]. 2024 Dec 13;12(24):3935. Available from: <https://www.mdpi.com/2227-7390/12/24/3935>

39.Nembrini S, König IR, Wright MN. The revival of the Gini importance? Valencia A, editor. *Bioinformatics* [Internet]. 2018 May 10;34(21):3711–8. Available from: <https://repository.publisso.de/resource/fri:6411640/data>

40.Elgart M, Lyons G, Romero-Brufau S, Kurniansyah N, Brody JA, Guo X, et al. Non-linear machine learning models incorporating SNPs and PRS improve polygenic prediction in diverse human populations. *Communications Biology* [Internet]. 2022 Aug 22 [cited 2023 Aug 28];5(1):1–12. Available from: <https://www.nature.com/articles/s42003-022-03812-z>

41.Wang H, Liang Q, Hancock JT, Khoshgoftaar TM. Feature selection strategies: a comparative analysis of SHAP-value and importance-based methods. *Journal of big data*. 2024 Mar 26;11(1).

42.Dhaarani R, Reddy MK. Progressing microbial genomics: Artificial intelligence and deep learning driven advances in genome analysis and therapeutics. *Intelligence-Based Medicine* [Internet]. 2025 Apr 23;11:100251. Available from: <https://www.sciencedirect.com/science/article/pii/S2666521225000559>

43.Divate M, Tyagi A, Richard DJ, Prasad PA, Gowda H, Nagaraj SH. Deep Learning-Based Pan-Cancer Classification Model Reveals Tissue-of-Origin Specific Gene Expression Signatures. *Cancers*. 2022 Feb 24;14(5):1185.

5 Supplementary Material




Figure 5: S1. Grid search performance curves for the top 10 hyperparameter configurations in region-level classification. Each line represents one of the top 10 Random Forest configurations, denoted as θ_0 through θ_9 , ranked by macro F1 score under 5-fold stratified cross-validation. Panels show macro F1 scores as a function of (A) number of estimators, (B) maximum tree depth, and (C) feature selection strategy ($max_features$). The red dot indicates the selected configuration (θ_7).




Figure 6: S2. Grid search performance curves for the top 10 hyperparameter configurations in country-level classification. Each line corresponds to a distinct hyperparameter configuration, θ_0 to θ_9 , outlining the top 10 performers ranked by macro F1 score across 5-fold stratified cross-validation. Subplots display macro F1 scores as functions of (A) number of estimators, (B) maximum tree depth, and (C) $max_features$ values. The red dot indicates the selected configuration (θ_3).




Figure 7: S3. Confusion matrices for region (left) and country (right) classification on the SMOTE-balanced training set.

Table 5: S4A. Per-Class Classification Metrics for Region Task (2019 Test Set)

Region Class	F1 Score	Precision	Recall
UK	0.900	1.000	0.819
S. Europe	0.844	0.826	0.864
M. East	0.894	0.927	0.864
N. Africa	0.850	0.773	0.944
Asia	0.480	0.400	0.600
C. Europe	0.667	0.524	0.917
S.S. Africa	0.000	0.000	0.000
C. America	0.471	0.333	0.800
N. Europe	0.667	0.556	0.833
N. America	0.400	0.250	1.000
Other Region	0.118	0.067	0.500

Table 6: S4B. Per-Class Classification Metrics for Country Task (2019 Test Set)

Country Class	F1 Score	Precision	Recall
UK	0.819	1.000	0.693
Other Country	0.758	0.766	0.750
Turkey	0.795	0.784	0.806
Spain	0.621	0.500	0.818
Egypt	0.500	0.429	0.600
Cyprus	0.615	0.471	0.889
Portugal	0.400	0.364	0.444
Greece	0.250	0.182	0.400
Morocco	0.167	0.111	0.333
Indonesia	0.286	0.182	0.667
France	0.250	0.154	0.667
Italy	0.400	0.250	1.000
India	0.000	0.000	0.000
Mexico	0.286	0.200	0.500
Ireland	0.143	0.077	1.000

Table 7: S4C. Overall and Per-Class Classification Metrics for Region Task (Training Set)

Overall Metrics			
Accuracy	0.995		
Macro F1	0.995		
Weighted F1	0.995		
Macro Precision	0.995		
Macro Recall	0.995		
Region Class	F1 Score	Precision	Recall
UK	0.996	0.996	0.996
S. Europe	0.992	0.994	0.990
M. East	0.995	0.996	0.994
N. Africa	0.996	0.996	0.997
Asia	0.997	0.998	0.997
Central Europe	0.996	0.997	0.996
S.S. Africa	0.993	0.993	0.993
C. America	0.994	0.995	0.994
N. Europe	0.994	0.990	0.997
N. America	0.995	0.994	0.995
Other Region	0.996	0.996	0.996

Table 8: S4D. Overall and Per-Class Classification Metrics for Country Task (Training Set)

Overall Metrics			
Accuracy	0.992		
Macro F1	0.926		
Weighted F1	0.992		
Macro Precision	0.926		
Macro Recall	0.926		
Country Class	F1 Score	Precision	Recall
UK	0.994	0.995	0.993
Turkey	0.992	0.992	0.992
Spain	0.992	0.990	0.994
Egypt	0.991	0.993	0.989
Cyprus	0.991	0.990	0.992
Portugal	0.994	0.993	0.994
Greece	0.991	0.992	0.991
Morocco	0.992	0.994	0.989
Indonesia	0.991	0.991	0.990
France	0.989	0.989	0.989
Italy	0.993	0.991	0.995
India	0.993	0.993	0.992
Mexico	0.993	0.992	0.994
Ireland	0.993	0.995	0.991

Table 9: S5A. Top 10 Predictive K-mers for Region Classification

Rank	Importance	K-mer Sequence
1	0.02407	TTTTTTTGCTTTTTGTAAATATCAATTTGTTATTTTTTGATTCAATTTGCTGGAGAGTATTTCAACAATTTTATTTTTTTT
2	0.02345	TTTTTGTGGGACCGCCGAAACATAACCGGACGGTGAGGGG
3	0.01507	TTTTTTAACGTCTTTGCCGAAACGCTGTTTTAGCACCGGTTAACCACAGCGTGGAATATAACTGGCTGAGGGAGGCGATCCCGTTTTTACAACAGGAA
4	0.01418	TTTTTGAGTTATTTCGTTTCGCAAGGTCTGACGGCGACGCCCTGACAACATCATAGTGTTAAATGCCATGAATCCTCCCGCCGGGATAATATTGTG
5	0.01411	TTTTTGAATCTTAAGTTTATCTTTCTGTTCTGCTCCTCTCGTCGTGTTCTTCTCTGCTGCTTTTCCGCTTTTTCGCGTTCTTACTTCGTGCTTCG
6	0.01369	TTTTTGCAACAGC
7	0.01326	TTTTTTCCGCAAAGCTACCGACCTGCTGAATGATTTTCACGCCATCAAACGCAGCATCGATACTATTTTCGCAGGGCATAGAAGCGCGTGTGCGCATTTGT
8	0.01287	TTTTTGCTATCC
9	0.01284	TTTTTTGACGGG
10	0.01275	TTTTTGATCGCTGA

Table 10: S5B. Top 10 Predictive K-mers for Country Classification

Rank	Importance	K-mer Sequence
1	0.02121	TTTTTTCTATCTGCCAACAAGTCAGCACGAGAACATTTTCGCCTCAATCAGGCATGATGCTGAATTTCTGAATCCCATAGCATCTGGCTGTTCTCCGGTAC
2	0.01342	TTTTTTGTGGTGCGTGAAGTGCGCCAGCTTATAAG
3	0.01245	TTTTTGTTCGACCTCCCTGGTCAGCATATCCTGTGCTTCAGATTTTACTGGCCGCTCTTCGGTTACCCCGGATTACCTTCGCCAGCCAGAAACCAGGTG
4	0.01214	TTTTTGCCCGAGCCCGGTACGAATGTTAATGGCGCAAATTTCTGGCGTCTCTTCTTCGCCGCTGCGCGCGCTTTCCACGGAAATAGCGCGCCCTGCCAG
5	0.01163	TTTTTGCTATG
6	0.01145	TTTTTGTTAACATTTAATATAATTATTATTAACCTCATGGACGCGTTAATGGCTAACTCATAATGGGTATTCAATAAGCTGCATTCTGTGATTGTTATCA
7	0.01140	TTTTTTGTCCGAGTGGCTTCGGCGTCCAGTTATCCTGACGAGCTTCTTATTTGCCGCTGCCGTTGCCGCTCTTTTGTAGCGGCAGCCTGGGCATTT
8	0.01119	TTTTTGTTTTTTCT
9	0.01105	TTTTTTCTCGCTC
10	0.01103	TTTTTTCCGCTGG